


Creating *reproducible, reusable, and scalable* bioinformatics workflows using Snakemake

Ebrahim Afyounian
Tampere - 09.09.19

What is a workflow?

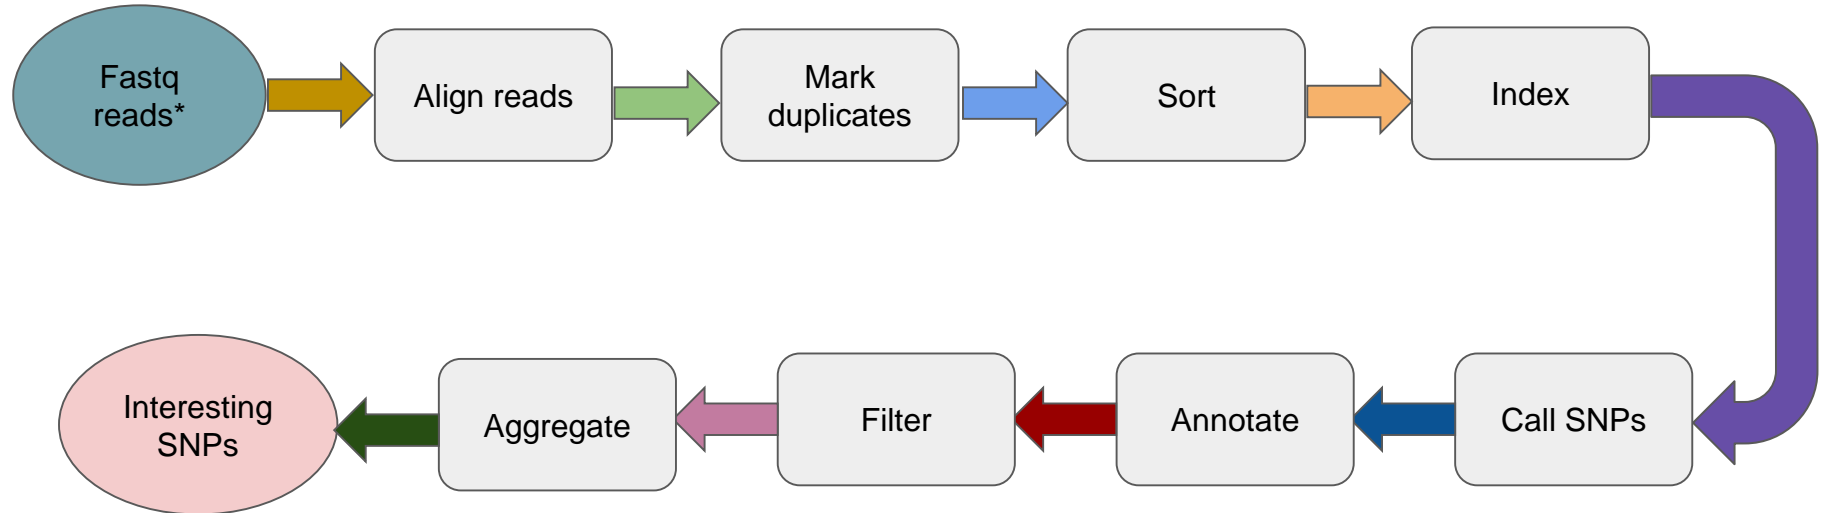


Workflow [\ 'wɜrk-, flō]. the **sequence of steps** involved in moving from the **beginning** to the **end** of a working process





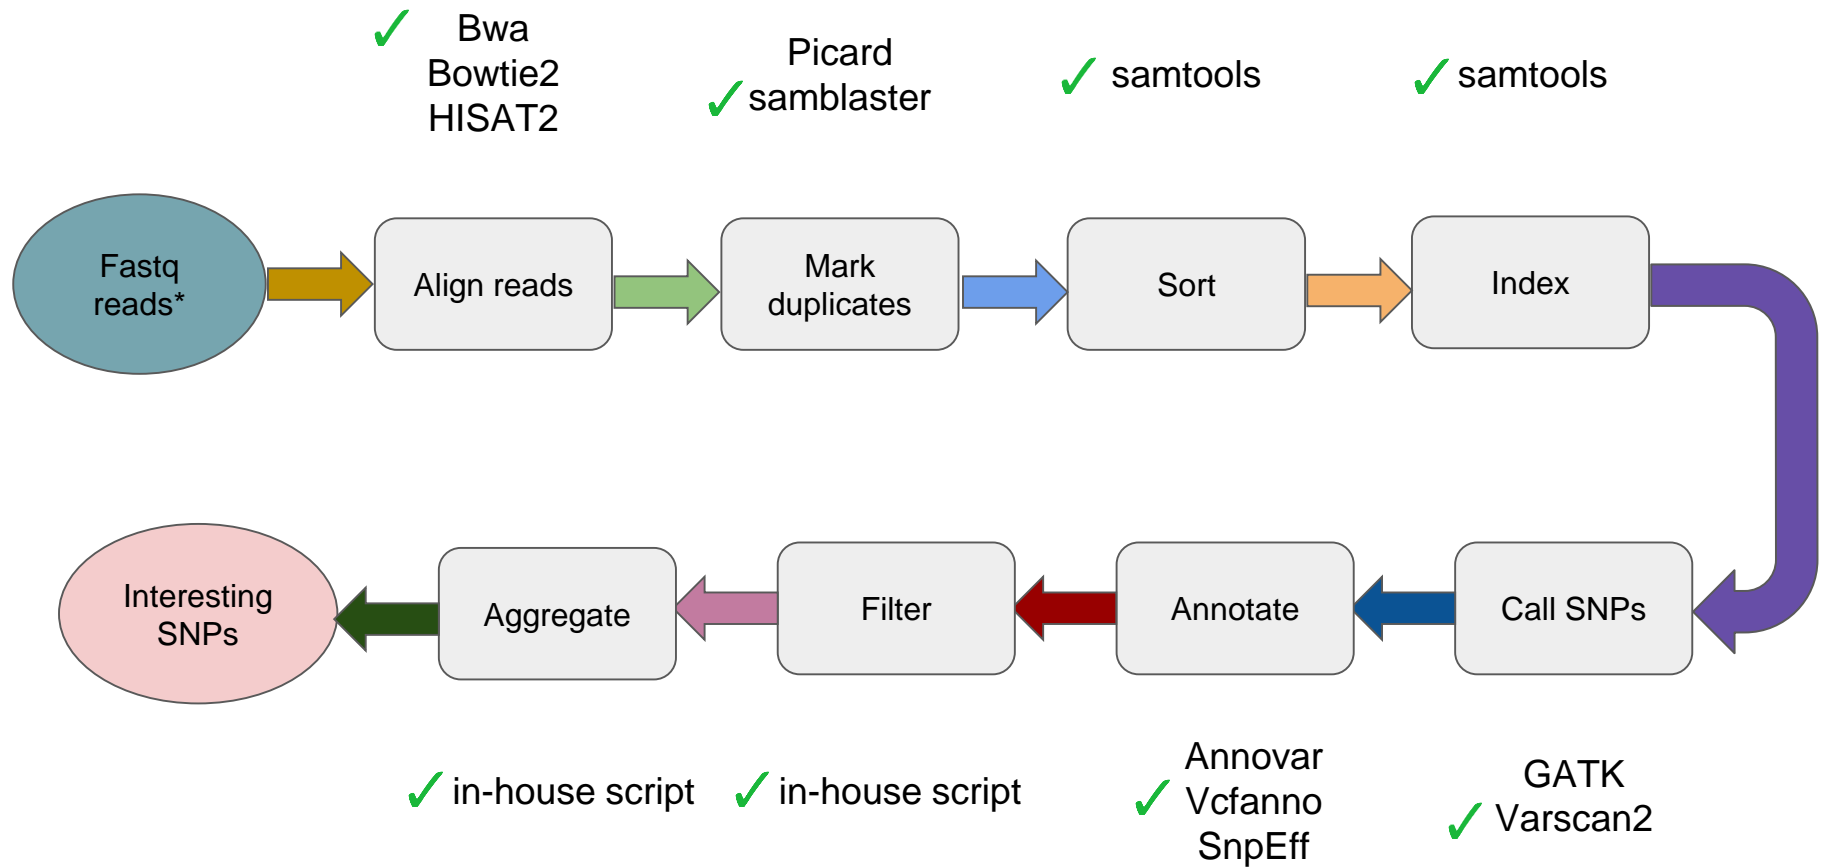
A mock example of a bioinformatics workflow



* Let's, for now, assume that we have done quality control (QC) over our fastq files, and we are happy with them.

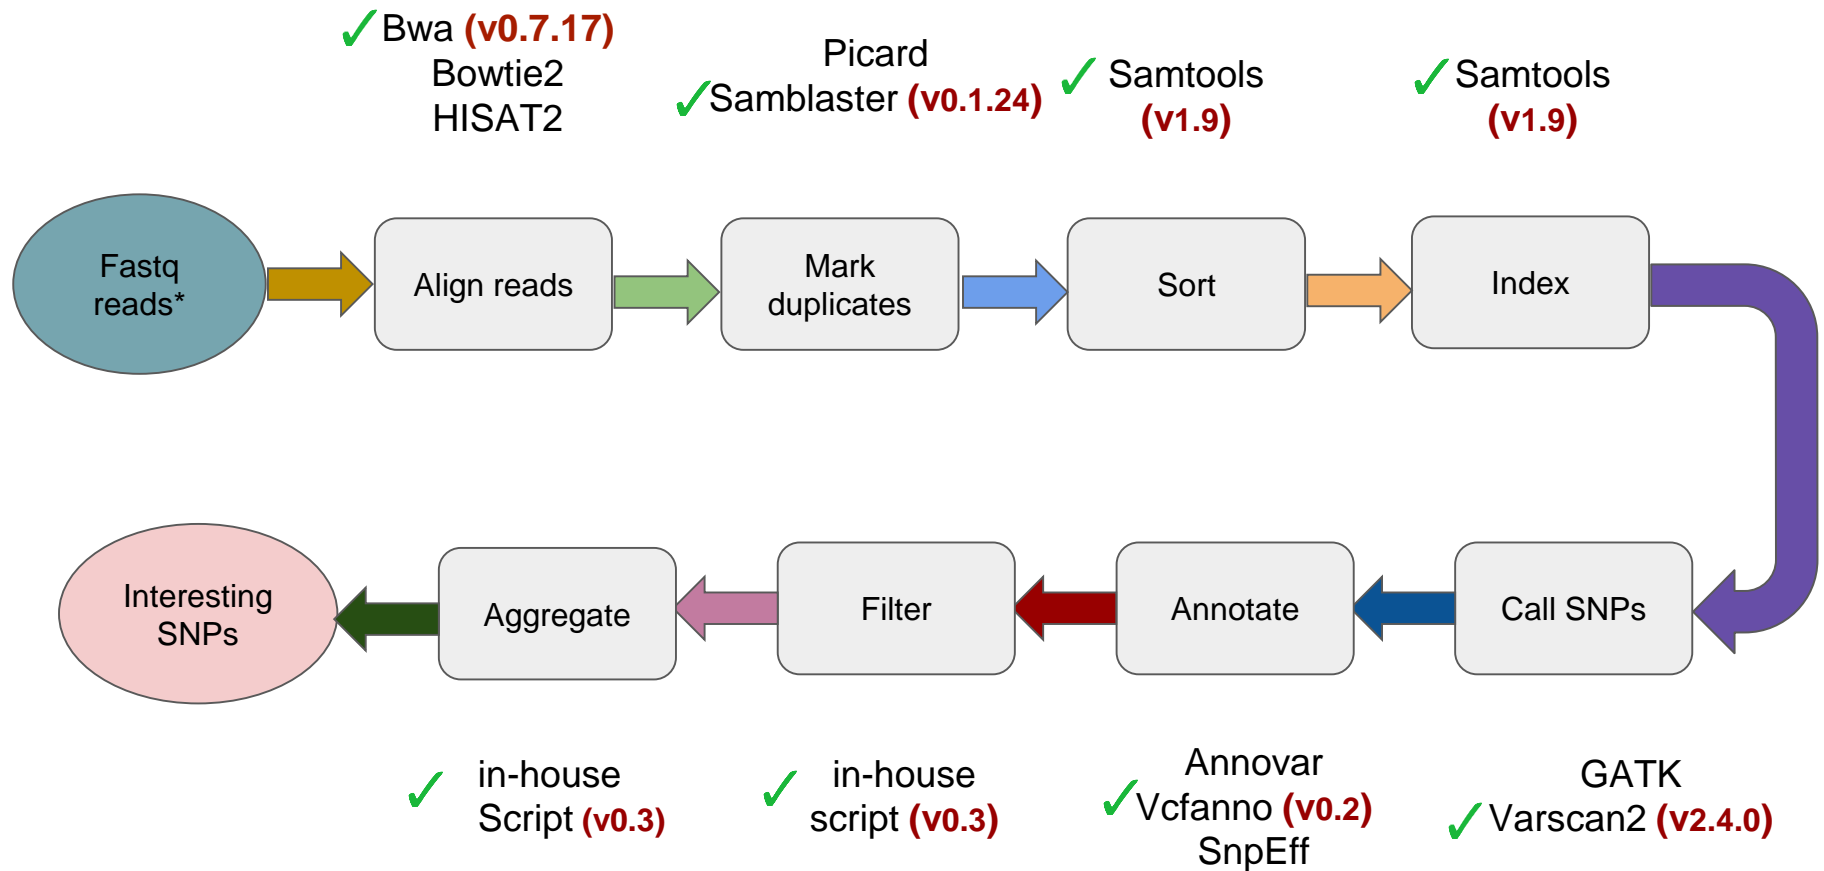


A mock example of a bioinformatics workflow



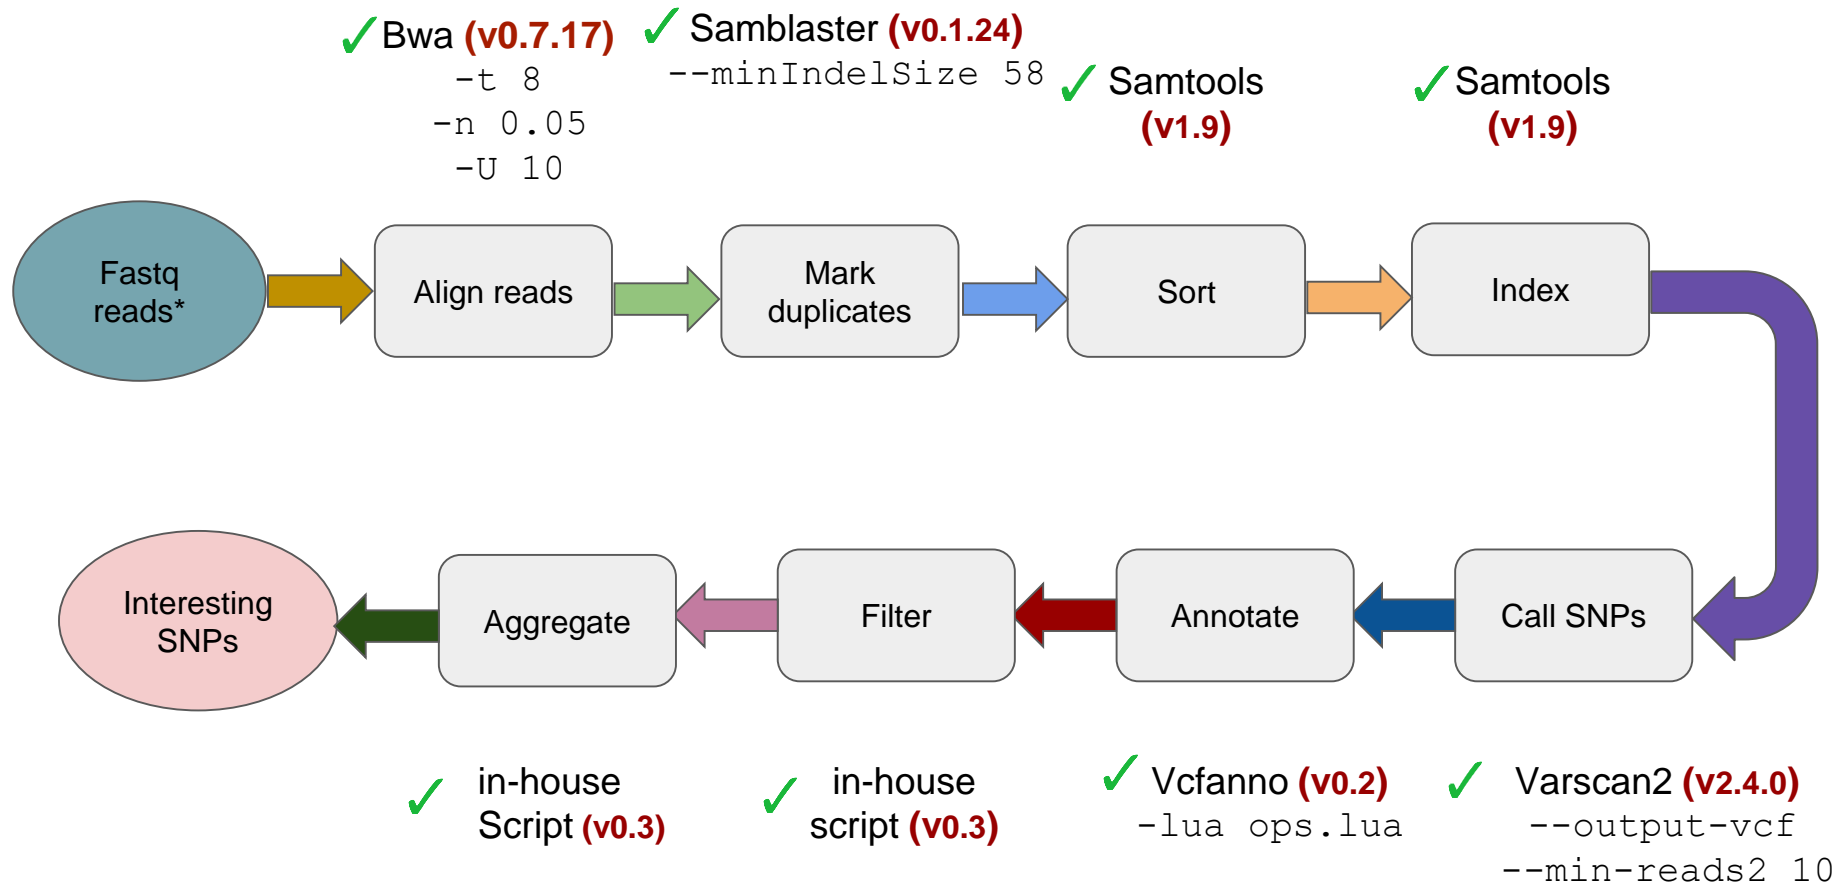


A mock example of a bioinformatics workflow





A mock example of a bioinformatics workflow



Scenario 1.



Let's imagine that our workflow works perfectly on one sample and:

- We receive 99 samples (198 fastq files).
- We are told that we have a limited amount of time (preferably ASAP) to find the *interesting SNPs* from these samples. In addition, we are required to provide our results with a report.



Option: Run our workflow in a ***for-loop*** and run one sample at a time.



Problem: what if something happens after processing x samples?

- Should we rerun every from scratch?
- Should we modify our for-loop so that it runs over only the unprocessed samples?



Problem: what if it will never finish by the deadline?



Option: Submit 99 jobs (1 per sample) to the job scheduler (e.g. to slurm)



Problem: what if an step in our workflow requires lots of *computational resources*. In such case, we need to request resources that are as big as the maximum amount required for that step, for the whole duration of the workflow's execution → We might end up **queueing for a long time** before our job starts. After, others have to queue until our jobs end!



Scenario 2.



*Let's imagine that we survived the 1st scenario **but**:*



Based on the results, our supervisor(s) decide(s) that: we should use stricter criteria for calling the SNPs (e.g. we need to change the parameter value for `-min-reads2` in `varscan2` from 10 to 20). **And while we are at it**, let's also add *gnomAD* annotation to our SNPs. **By the way**, we need the results for **tomorrow afternoon's meeting** with the people from the clinic [... **and pliiis** *bring some sweets to the meeting ...*].



Question: Should we update the criteria and rerun our entire workflow from scratch?



We know that simply our workflow won't finish in less than a day!



Question: Or should we modify our workflow so that it runs only from `call SNPs` step and possibly save time and the resources?



Possible, but then we have broken our workflow into smaller pieces i.e. we have one workflow for scenario 1 and another for scenario 2.



Bioinformatics workflow management systems (WMS)

Snakemake 

nextflow



Anduril is from *Systems Biology Laboratory*, University of Helsinki

Snakemake—a scalable bioinformatics workflow engine FREE


Johannes Köster, Sven Rahmann

Bioinformatics, Volume 34, Issue 20, 15 October 2018, Page 3600,

<https://doi.org/10.1093/bioinformatics/bty350>

Published: 16 May 2018



- A *text-based* workflow management system introduced in 2012
- Inherits many of its features from GNU `make`.
- Implements a domain-specific language (an extension of Python ) with which we define/formalize the steps in our data analysis as a workflow
- A workflow is composed of **rules** i.e. Each rule is one step in the data analysis / workflow
- Each rule has ***input***, ***output***, and ***the computation that turns the input to output***
- Having rules makes the workflow **modular**.
 - + It can be **reused** in another workflow
 - + We can request computational resources for each rule **separately**
- Checks the dependencies among the rules (e.g. ***rule 2*** needs ***rule 1***'s result as input) and infers a dependency graph and execution order
- Decides what steps need to be done in the analysis ensuring that we do/re-do only what is needed
- If possible, rules can be executed parallelly.

Basic anatomy of a rule

```
rule bwa_map:
    input:
        "data/genome.fa",
        "data/samples/A.fastq"
    output:
        "mapped_reads/A.bam"
    shell:
        "bwa mem -t8 {input} | samtools view -O BAM @7 - > {output}"
```

- input, output, and shell are called **directives**
- There are few more very useful directive:
 - script: used to run scripts such as Python and R
 - conda: used to manage environments and packages
 - thread: used to specify number of threads
 - log: used to specify where the log info is written



What happens under the hood?

```
rule bwa_map:
    input:
        "data/genome.fa",
        "data/samples/A.fastq"
    output:
        "mapped_reads/A.bam"
    shell:
        "bwa mem -t8 {input} | samtools view -O BAM @7 - > {output}"
```

```
"bwa mem -t8 data/genome.fa data/samples/A.fastq | samtools view -O BAM @7 - > mapped_reads/A.bam"
```

Some issues to be addressed:

- Not scalable

Hard-coding the sample name(s)

- Not necessarily reusable:

It assumes that `bwa` and `samtools` are installed on the system (which versions we do not know)?



Addressing some of the issues (1)



```
rule bwa_map:
    input:
        "data/genome.fa",
        "data/samples/A.fastq"
    output:
        "mapped_reads/A.bam"
    shell:
        "bwa mem -t8 {input} | samtools view -O BAM @7 - > {output}"
```



```
rule bwa_map:
    input:
        "data/genome.fa",
        "data/samples/{sample}.fastq"
    output:
        "mapped_reads/{sample}.bam"
    shell:
        "bwa mem -t8 {input} | samtools view -O BAM @7 - > {output}"
```

Better! But not quite!

No hard-coded samples

Reference is hard-coded

Addressing some of the issues (2)



```
rule bwa_map:
    input:
        "data/genome.fa",
        "data/samples/{sample}.fastq"
    output:
        "mapped_reads/{sample}.bam"
    conda:
        "environments/alignment.yaml"
    shell:
        "bwa mem -t8 {input} | samtools view -O BAM @7 - > {output}"
```



```
channels:
    - conda-forge
    - bioconda
dependencies:
    - bwa=0.7.17
    - samtools=1.9
```

- No hard-coded samples
- Uses `conda*` to create an environment for the rule
- Installs required tools in the environment.



Information about where the sample(s) are still missing.

* Conda is an open source **environment management** and **package management** system



Running a Snakemake workflow

- Snakemake workflow can be run on:
Desktop machine; Server / Cluster; Cloud
- When running a workflow, we need to specify a **target** i.e. the *output(s)* we want to have.
- If a **target** is not provided, Snakemake runs the **first rule** by default as well as any other rule needed for successful completion of the first rule.
- In practice, the first rule is the “rule all”.

```
$ snakemake --dag | dot -Tsvg > dag.svg ## creates a directed acyclic graph (DAG) visualization
$ snakemake --dryrun mapped_reads/A.bam ## Snakemake tells what is going to happen
$ snakemake --use-conda mapped_reads/A.bam ## use --use-conda if we have conda directive
$ snakemake --use-conda ## runs starting from the first rule since no target mentioned
$ sankemake --rerun-incomplete ## only rerun incomplete rules

## more complex with more options. Sending jobs to the job scheduler
## (i.e. slurm). Not the best way though! since it request same resources
## for each rule (i.e. possibility for over-calculation of the resources).
$ snakemake --cluster "sbatch --cpus-per-task=1 --nodes=1 --mem-per-cpu=16G --
partition=parallel,test,normal --time=01:00:00 --parsable" --jobs 50 --latency-wait 120 --nolock

## Should fixe the above problem. We need to set the params directive
## for each rule. I have not tested this! So not 100% sure if this works.
## this is cumbersome. One can define profiles that are YAML files for this.
$ snakemake --cluster "sbatch --time {params.time} --mem-per-cpu {params.memory}"
```



An example workflow

```
configfile: "config.yaml"

rule all:
  input:
    ["sorted_reads/%s.bam" %sample for sample, path in config['samples'].items()]
  shell:
    "echo done!"

rule bwa_map:
  input:
    ref_gen=config["reference_genome"],
    sample=lambda wildcards: config["samples"][wildcards.sample]
  output:
    temp("mapped_reads/{sample}.bam")
  conda:
    "environments/alignment.yaml"
  shell:
    "bwa mem -t8 {input.ref_gen} {input.sample} | samtools view -O BAM @7 - > {output}"

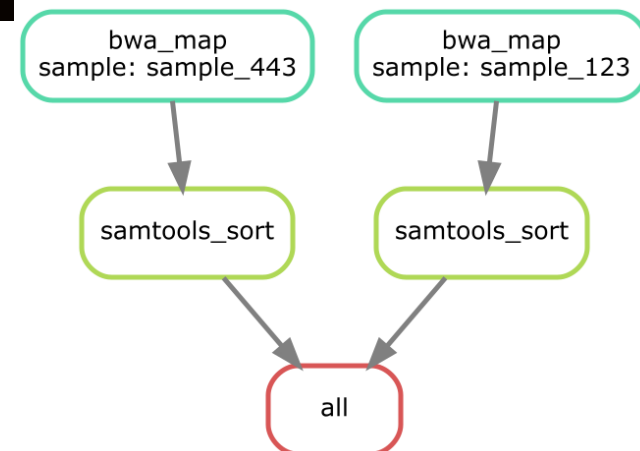
rule samtools_sort:
  input:
    "mapped_reads/{sample}.bam"
  output:
    "sorted_reads/{sample}.bam"
  conda:
    "environments/alignment.yaml"
  shell:
    protected("samtools sort -T sorted_reads/{wildcards.sample}"
    "-O bam {input} > {output}")
```

Config file

```
samples:
  sample_123: data/sample_123.fastq
  sample_443: data/sample_443.fastq

reference_genome:
  data/references/genome.fa
```

Dependency graph



Dry run output



```
Building DAG of jobs...
```

```
Job counts:
```

count	jobs
1	all
2	bwa_map
2	samtools_sort
5	

```
[Sun Sep  8 22:57:15 2019]
```

```
rule bwa_map:
```

```
  input: data/references/genome.fa, data/sample_443.fastq
  output: mapped_reads/sample_443.bam
  jobid: 4
  wildcards: sample=sample_443
```

```
bwa mem -t8 data/references/genome.fa data/sample_443.fastq | samtools view -O BAM @7 - > mapped_reads/sample_443.bam
```

```
[Sun Sep  8 22:57:15 2019]
```

```
rule bwa_map:
```

```
  input: data/references/genome.fa, data/sample_123.fastq
  output: mapped_reads/sample_123.bam
  jobid: 3
  wildcards: sample=sample_123
```

```
bwa mem -t8 data/references/genome.fa data/sample_123.fastq | samtools view -O BAM @7 - > mapped_reads/sample_123.bam
```

```
[Sun Sep  8 22:57:15 2019]
```

```
rule samtools_sort:
```

```
  input: mapped_reads/sample_443.bam
  output: sorted_reads/sample_443.bam
  jobid: 2
  wildcards: sample=sample_443
```

```
samtools sort -T sorted_reads/sample_443-O bam mapped_reads/sample_443.bam > sorted_reads/sample_443.bam
```

```
[Sun Sep  8 22:57:15 2019]
```

```
rule samtools_sort:
```

```
  input: mapped_reads/sample_123.bam
  output: sorted_reads/sample_123.bam
  jobid: 1
  wildcards: sample=sample_123
```

```
samtools sort -T sorted_reads/sample_123-O bam mapped_reads/sample_123.bam > sorted_reads/sample_123.bam
```

```
[Sun Sep  8 22:57:15 2019]
```


```
rule all:
```

```
  input: sorted_reads/sample_123.bam, sorted_reads/sample_443.bam
  jobid: 0
```




Few of the features not presented here

- Can create HTML reports showing e.g. the rules' run time, files creation time, even results / visualizations from the rules...

 `--report`


- Specify that a rule output should be *piped* into another rule preventing writing to disk

 `pipe()`

- `resources` directive can be used to set the resources needed to run a job
- Ability to run in containers (e.g. Docker, Singularity)
- And many more ...
- **Upcoming feature(s):**
 - Integration with ***Jupyter notebook*** (i.e. jupyter directive)

More resources



- Snakemake homepage:
<https://snakemake.readthedocs.io>
 - Snakemake tutorial:
<https://snakemake.readthedocs.io/en/stable/tutorial/tutorial.html>
 - Bioinformatics workflows using Snakemake (e.g. DNA-seq GATK variant calling, Single-cell RNA-seq analysis):
<https://github.com/snakemake-workflows/docs>
 - Collection of reusable wrappers allowing use of popular tools (such as `fatsqc`, `samtools`, `bwa`, etc.) from Snakemake rules and workflows:
<https://snakemake-wrappers.readthedocs.io/en/stable/>
 - Live demo created by the Johannes Köster:
<https://www.katacoda.com/johanneskoester/scenarios/snakemake-intro>
-  No need to use your real email address to run the live demo. Just type a dummy email address and password in the sign-up form and continue.
- Nice presentation by Johannes Köster:
<https://youtu.be/hPrXcUUp70Yp70Y>